

# VizAPI: Visualizing Interactions between Java Libraries and Clients

Sruthi Venkatanarayanan, Patrick Lam  
University of Waterloo  
Waterloo, ON, Canada  
{s42venkat,patrick.lam}@uwaterloo.ca

Jens Dietrich, Craig Anslow  
Victoria University of Wellington  
Wellington, New Zealand  
{jens.dietrich,craig.anslow}@vuw.ac.nz

**Abstract**—Software projects make use of libraries extensively. Libraries make available intended API surfaces—sets of exposed library interfaces that library developers expect clients to use. However, in practice, clients only use small fractions of intended API surfaces of libraries. We have implemented the VizAPI tool, which shows a visualization that includes both static and dynamic interactions between clients, the libraries they use, and those libraries’ transitive dependencies (all written in Java). We then present some usage scenarios of VizAPI. One application, by client developers, is to answer a query about upstream code: will their code be affected by breaking changes in library APIs? Additionally, library developers can use VizAPI to find out about downstream code: which APIs in their source code are commonly used by clients?

**Index Terms**—static program analysis, dynamic program analysis, API usage, software evolution, software maintenance

## I. INTRODUCTION

Virtually all modern software projects use libraries, driven in part by the ease of use of open source component repositories such as Maven and npm. Library developers design Application Programming Interfaces, or APIs, for their libraries, and clients invoke these APIs. APIs typically provide clients with methods that can be invoked, fields that can be accessed, classes that can be instantiated, and annotations that can be used. We use the term “API surface” to denote the APIs that a library makes available to other code artifacts (clients and other libraries). Myers and Stylos [18], and many others, have advocated for the importance of easy-to-use and maintainable API surfaces.

The number of dependencies used by modern software has exploded, and so has their complexity [2], [11]: deeper, transitive dependencies are now common, components are upgraded more frequently, and developers increasingly struggle to deal with issues arising from those changes. Issues include: (1) dealing with conflicting versions of the same component (also known as dependency hell) and dealing with supply chain vulnerabilities of deep dependencies (often notified by bots creating pull requests); (2) new issues around security and resilience of the software supply chains, e.g. problems with changes to commodity components (as in the infamous left-pad incident [4]) and novel attack patterns like typosquatting; and, (3) the use of unnecessary, bloated, and trivial dependencies [1], [20].

This work was partially supported by Canada’s Natural Science and Engineering Research Council.

The benefits of using libraries (e.g. the ease of including functionality that one is not responsible for maintaining) are thus offset by the issues mentioned above. Let’s consider one other issue: breaking changes. *Potentially* breaking changes in library APIs are common [6], [19]. However, any library change is only potentially breaking; does it *actually* break any particular client? Only if a client uses a specific component API with an incompatible change.

In all programming environments that we are familiar with, a client developer’s decision to use a particular API (e.g. library method  $m()$ ) is a local decision, taken by the developer as they build their system. To our knowledge, there are no existing systems that help developers visualize which APIs are used throughout an entire project. We claim that such a visualization has potential applications for both client and library developers.

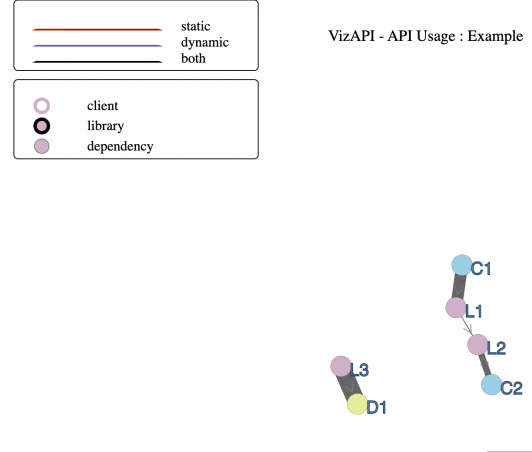


Fig. 1. An Example VizAPI Application

Figure 1 illustrates a possible VizAPI usage scenario, from the perspective of a client developer. Consider a client  $C$  (blue nodes) and a library  $L$  (purple nodes), in the context of plain Java. Library  $L$  has packages  $L_1$ ,  $L_2$ , and  $L_3$ .  $C$  calls into  $L_1$  and  $L_2$ . Internally, within  $L$ ,  $L_1$  and  $L_2$  call into each other, but not into  $L_3$ . The VizAPI result, with no edges from  $C$  directly to  $L_3$ , allows a developer to conclude that breaking changes in  $L_3$  will not affect  $C$ . Also, if only  $L_3$  uses an external dependency  $D$  (yellow node), then we know that  $C$  will not need  $D$  to be on its classpath.

From the library developer side, under plain Java (i.e. no runtime containers) and considering reflection, the potential API surface of any component is huge. Essentially: every method can be called, and every field can be read and written. Even considering only the published API surface (methods, fields, classes, and annotations with the correct visibility modifiers), libraries' API surfaces still often have hundreds to thousands of members. The breadth of the API surface is a liability with respect to continued maintenance of the library; many developers aspire to avoid breaking changes by preserving, whenever possible, library behaviour that is depended on by clients. Knowing that few clients use a particular API would be valuable.

On the other hand, we would expect (and have verified in other work) that each client uses only a small portion of each of its dependencies' API surfaces. Consider breaking changes again. GitHub provides the Dependabot tool [17], which monitors for upstream changes and automatically proposes pull requests to update dependency versions. That tool may well pull in breaking changes. However, we hypothesize that, most of the time, most breaking changes will not affect most clients; it is useful for clients to know whether they are using the parts of the API surface that are subject to a particular breaking change. A client with broad dependencies on a library (uses a larger fraction of its API surface) is more likely to be affected by its changes than a client with narrow dependencies (smaller fraction). A narrow library dependency would also suggest that it would be easier to swap the library for a functionally similar replacement.

Additionally, as researchers, we would like to understand how library APIs are used by clients more generally. Zhong and Mei [23] investigated API usages in a dataset of 7 experimental subjects (clients) and the libraries that they depend on. They found that clients use less than 10% of the declared APIs in libraries. Our visualization allows developers and researchers to visualize distribution information about how different parts of clients use different parts of libraries.

This paper presents the VizAPI tool, which shows visualization overviews showing API usages—from clients to libraries, but also between libraries (including transitive dependencies). VizAPI incorporates information from static and dynamic analyses. We have made VizAPI publicly available, although it is still in development:

<https://github.com/SruthiVenkat/api-visualization-tool>

Our contributions include:

- an implementation of the VizAPI tool, which presents a visualization of API usage information; VizAPI collects static information and instruments Java code and collect dynamic instrumentation information about API uses in practice and presents it as a d3 visualization (Section III-A–III-B);
- a discussion of VizAPI usage scenarios (Section III-C) based on a collection of 11 libraries and 38 clients.

## II. RELATED WORK

A representative tool from the software visualization literature is CodeSurveyor, by Hawes et al [9]. This tool visualizes large codebases using cartographic maps as an analogy. While it incorporates dependency information into the layout of the map, VizAPI differs from CodeSurveyor in that VizAPI focusses on usage relationships between different modules—primarily API invocations—using test cases and static analysis to identify relationships between clients and libraries, rather than investigating a particular system, as CodeSurveyor does. Earlier work in the same vein includes the software cartography project by Kuhn et al [12] and software terrain maps by DeLine [5].

Hejderup and Gousios [10] explore a question which is central to our approach—how well do client tests exercise their dependencies' libraries? That is, we rely on client test suites exercising enough of the dependencies to get valid results from our analyses. Their conclusion is that a combination of static and dynamic analysis of the client has some chance of detecting breaking changes in its dependencies, and we accordingly use static analysis to supplement our dynamic results.

Thummalapenta and Xie [21] presented the related SpotWeb tool, which identifies framework hotspots (APIs that are often used) and coldspots (APIs that are never used); they do not consider framework APIs that are used but not intended to be. Hotspots and coldspots are, however, related to our investigations about client use of APIs; part of our study could be seen as investigating the prevalence of hotspots on our benchmark suite. Their notion of API usage is similar to ours, but they perform a static search to identify uses, while we additionally dynamically observe test executions. They also identify the top  $N$  percent of used APIs as hotspots, and unused APIs as coldspots. Viljamaa [22] also aimed to find hotspots but used concept analysis to do so.

Our overall goal is to help both client and library developers understand client uses of library code. Clients benefit from sharpened warnings about unsafe upgrades, knowledge that some upgrades are safe, and having reduced attack surfaces. Library upgrades have been investigated by many researchers, including Lam et al [14] and Kura et al [13]. Kura et al found that most software had outdated dependencies, and that software developers had negative feelings about being required to constantly upgrade their libraries. Being able to visualize dependencies may help developers prioritize required upgrades as low-effort or high-effort. Foo et al [8] proposed a static analysis which detected safe upgrades, but could only certify safety for 10% of upgrades; our combined static and dynamic approach presents the developer with more information and enables more upgrades.

## III. VIZAPI

We next describe the design of VizAPI, including details on how we collect information and format it for the d3js visualization library. We also present two case studies showing potential uses of VizAPI.

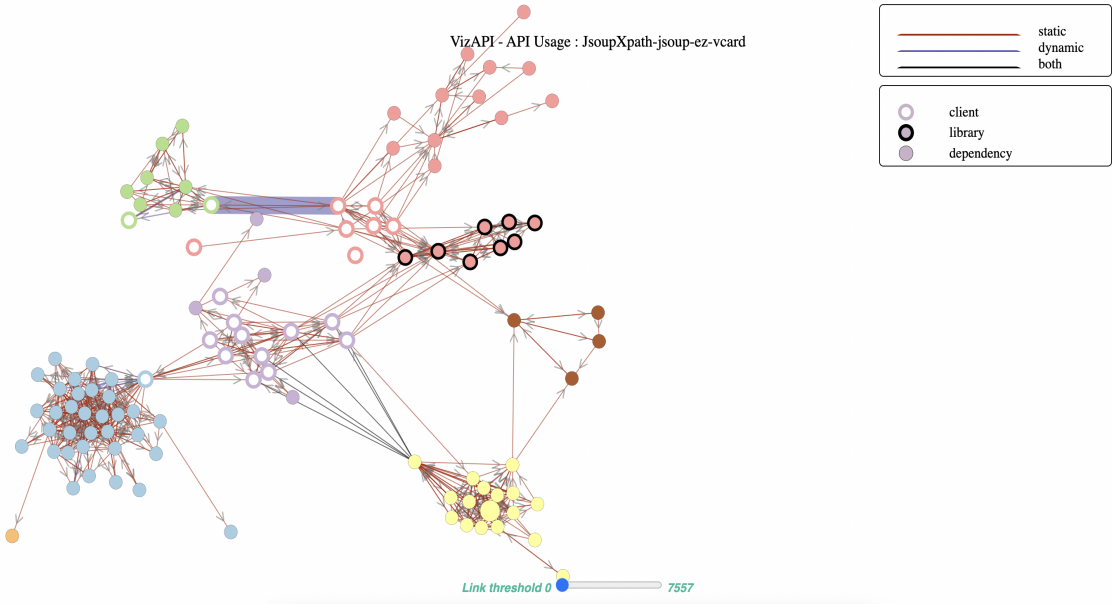


Fig. 2. Usage Scenario 1: jsoup (library), JsoupXPath (client), ez-vcard (client)

### A. Design

Our goal is to visualize interactions between different software components—between clients and libraries, and between libraries and other libraries. We now explain how we collect static and dynamic information about software behaviour. Figure 3 summarizes our instrumentation and data capture workflow.

To identify interactions across the client/library boundaries, we first inspect JAR files of each software component to obtain a list of classes for every component. We then associate classes and their members to components based on these lists.

*a) Static information:* We perform a static analysis to record API uses. Using Javassist [3], we identify type references, which includes references at call sites, fields, annotations for classes, methods and fields, method parameters and casts. We also identify subtypes, determining possible interactions across client/library boundaries and library/library boundaries. Javassist uses a call graph produced by class hierarchy analysis to resolve method calls.

*b) Dynamic information:* We collect dynamic data by running client test suites under instrumentation. The instrumentation records API uses which cross client/library boundaries, as well as library/library boundaries for libraries that are transitively used. We also use Javassist to perform this instrumentation and then use the build system of each project (Maven or Gradle) to run its tests.

At every invoke instruction in every loaded method which transfers control between the client and the library, we add code to record that invoke by incrementing a counter. We handle both static and virtual (including special, virtual, interface, and dynamic) calls. Crossing the client/library boundary includes callbacks from the library to the client as well as conventional calls from the client to the library.

We also record field accesses (direct and reflective), dynamic proxies and reflective calls, Java annotations, implementations, instantiations, and casts.

### B. Visualization System

Once we have generated data, we use a modified version of the d3graph<sup>1</sup> library in Python to generate a d3js<sup>2</sup> visualization.

VizAPI graphs are force-directed graphs based on the frequency of interactions between different software components. Each software component can be one of: a client; or a dependency. Both libraries and dependencies refer to external libraries imported by components. We refer to components used by clients as libraries, and components used by libraries as dependencies.

The graph in Figure 2 is an example of a graph produced by VizAPI. Each node is a set of one or more packages (or classes or methods) that belong to the same JAR. Clients are represented by nodes with white interiors; libraries by nodes with filled interiors and black borders; and dependencies by nodes with filled interiors and normal borders. We coalesce nodes if they originate from the same JAR file and have the same incoming and outgoing edges. Each edge is directed from the source package(s) to the target package(s) and represents an interaction (invocations, fields, annotations, subtyping) between packages. We run a clustering algorithm and use its output to colour nodes based on the cluster that they belong to, so a colour could include nodes from the same or different JARs. Hovering on a node shows the list of packages and the JAR that they belong to, formatted as “jar : <space separated list of packages>”.

<sup>1</sup><https://pypi.org/project/d3graph/>

<sup>2</sup><https://d3js.org/>




Fig. 3. Our instrumentation workflow. We modify existing project infrastructure to instrument clients and to run their test suites, producing raw data, which we process with our Python scripts to create D3 visualizations.

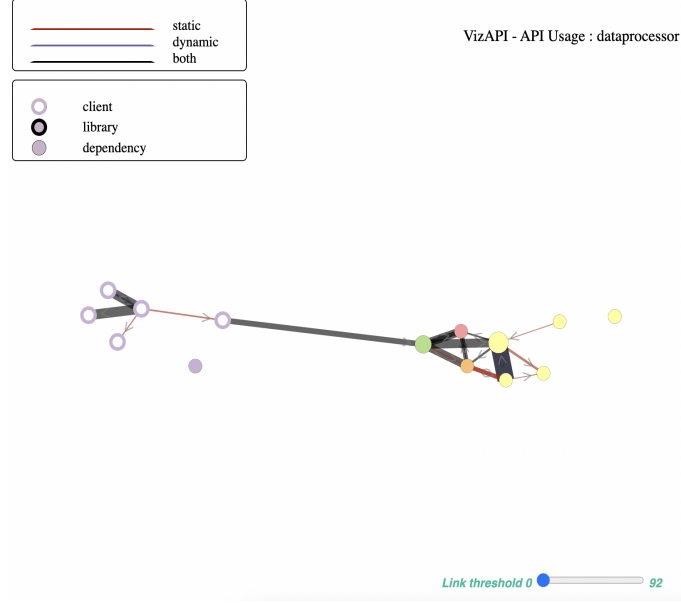


Fig. 4. Usage Scenario 2: dataprocessor (client)

### C. Case Study

We conducted a pilot study of VizAPI. We have generated data from libraries from a subset of the Duets benchmarks [7] combined with a selection of popular Maven repositories in different categories such as logging and json parsing. Our study included 10 libraries and selected clients of these libraries (potentially overlapping), for a total of 85 projects. We have made our data publicly available at

<https://www.dropbox.com/s/puopek9t9nbm8ip/apis-data.zip?dl=0>

We chose clients partly from popular Github repositories and partly from Duets. We have collected both static and dynamic data for these projects, and we are in a position to generate graphs for combinations of clients and libraries in these projects. We present two usage scenarios below; graphs for our usage scenarios can be found at:

<https://sruthivenkat.github.io/VizAPI-graph/>

a) *Usage Scenario 1: jsoup*: Now, we walk through an example usage scenario of VizAPI. The graph in Figure 2

represents static and dynamic interactions of 2 clients with the jsoup<sup>3</sup> library. Our clients are JsoupXpath<sup>4</sup> and ez-vcard<sup>5</sup>.

We can start our exploration with the cluster of pink nodes. Many of these nodes belong to either JsoupXpath or jsoup. When we hover over them, the hover hints tell us that the client JsoupXpath calls directly into `org.jsoup.nodes` and `org.jsoup.select`. Notably, and as we might expect, we can see that `org.jsoup.helper` and `org.jsoup.internal` aren't called directly by JsoupXpath. This would mean that breaking changes in `org.jsoup.helper` or `org.jsoup.internal` wouldn't directly affect JsoupXpath<sup>6</sup>

Similarly, ez-vcard, which belongs to the purple cluster in Figure 2, directly calls into `org.jsoup`. ez-vcard also

<sup>3</sup><https://github.com/jhy/jsoup>

<sup>4</sup><https://github.com/zhegexiaohuoz/JsoupXpath>

<sup>5</sup><https://github.com/mangstadt/ez-vcard>

<sup>6</sup>As a specific example, the retraction of an internal jsoup API would not break this client. Behavioural changes that are directly passed through to the external API, e.g. through delegation, can still break clients, but we can consider those to be changes in the external API.

calls into `jackson-core`<sup>7</sup> and `jackson-databind`<sup>8</sup>, which are very tightly coupled amongst their own packages and with each other. This could mean that breaking changes in these libraries can propagate to many of their own packages.

b) *Usage Scenario 2: dataprocessor*: Figure 4 presents a second usage scenario. Here, we focus on a client, `dataprocessor`<sup>9</sup>. This client uses the `fastjson`<sup>10</sup> library. Our visualization shows calls only from `dataprocessor` package `com.github.dataprocessor.slice`, which is the orange client node (identity of the package available from clicking on the orange node) to the package `com.alibaba.fastjson`. No other parts of `dataprocessor` use `fastjson`. This means that when `dataprocessor` needs to upgrade its `fastjson` version, the `dataprocessor` developers only need to inspect the source code in the `com.github.dataprocessor.slice` package.

Note also the disconnected nodes in Figure 4. These are all packages of `fastjson` that are unused by `dataprocessor`: any breaking changes in these packages definitely do not directly affect `dataprocessor`, and are less likely to affect it overall than packages that are directly used.

#### IV. DISCUSSION

Our goal when developing VizAPI was to enable 1) library developers to make better decisions about pruning or modifying unused APIs and to refactor their libraries; and 2) client developers to make better decisions about library upgrades and breaking changes.

Note that client tests may not adequately represent actual client behaviours; however, our use of both static and dynamic information addresses this issue. Specifically, because we use class hierarchy analysis for our static analysis, our visualization will present all possible static calls—possibly too many. That is, the main hazard with static analysis is that our visualization may include more static edges than are actually possible. Some of those edges could be ruled out by a more precise call graph. Reflection aside, no static edges are missing (our approach is “soundy [15]” with respect to static information). On the other hand, dynamic edges have actually been observed on some execution; better tests could yield more dynamic edges. But even if a dynamic edge is missing, there will be a static edge if the behaviour is possible.

Nevertheless, this preliminary work has presented two usage scenarios which we believe promise to be useful for both client and library developers. Consistent with the Call for Papers for the NIER track, we have not yet carried out a formal evaluation of our VizAPI tool. We intend to carry out further evaluation of our tool following the techniques described by Merino et al [16]; in particular, we aspire to perform experiments to establish the effectiveness of VizAPI, where we ask users to perform software understanding and maintenance tasks that would benefit from our tool.

<sup>7</sup><https://github.com/FasterXML/jackson-core>

<sup>8</sup><https://github.com/FasterXML/jackson-databind>

<sup>9</sup><https://github.com/dadiyang/dataprocessor>

<sup>10</sup><https://github.com/alibaba/fastjson>

#### REFERENCES

- [1] Rabe Abdalkareem, Olivier Nourry, Sultan Wehaibi, Suhaib Mujahid, and Emad Shihab. Why do developers use trivial packages? an empirical case study on npm. In *FSE*, pages 385–395, 2017.
- [2] Amine Benelallam, Nicolas Harrand, César Soto-Valero, Benoit Baudry, and Olivier Barais. The Maven dependency graph: a temporal graph-based representation of Maven Central. In *MSR*, pages 344–348. IEEE, 2019.
- [3] Shigeru Chiba. Load-time structural reflection in Java. In *ECOOP*, volume 1850 of *LNCS*, pages 313–336. Springer Verlag, 2000.
- [4] Keith Collins. How one programmer broke the internet by deleting a tiny piece of code. *Quartz*, March 2016. <https://qz.com/646467/how-one-programmer-broke-the-internet-by-deleting-a-tiny-piece-of-code/>. Accessed 19 October 2021.
- [5] Robert DeLine. Staying oriented with software terrain maps. In *Proceedings of the Workshop on Visual Languages and Computation*, 2005.
- [6] Jens Dietrich, Kamil Jezek, and Premek Brada. Broken promises: An empirical study into evolution problems in Java programs caused by library upgrades. In *WCRE*, pages 64–73. IEEE, 2014.
- [7] Thomas Durieux, César Soto-Valero, and Benoit Baudry. DUETS: A dataset of reproducible pairs of java library-clients. In *MSR*, 2021.
- [8] Darius Foo, Hendy Chua, Jason Yeo, Ming Yi Ang, and Asankhaya Sharma. Efficient static checking of library updates. In *FSE*, page 791–796, New York, NY, USA, 2018. Association for Computing Machinery. doi:10.1145/3236024.3275535.
- [9] Nathan Hawes, Stuart Marshall, and Craig Anslow. CodeSurveyor: Mapping large-space software to aid in code comprehension. In *VISSOFT*, Bremen, Germany, 2015.
- [10] Joseph Hejderup and Georgios Gousios. Can we trust tests to automate dependency updates? A case study of Java projects. *J. Syst. Softw.*, 183:111097, 2022. doi:10.1016/j.jss.2021.111097.
- [11] Riivo Kikas, Georgios Gousios, Marlon Dumas, and Dietmar Pfahl. Structure and evolution of package dependency networks. In *MSR*, pages 102–112. IEEE, 2017.
- [12] Adrian Kuhn, David Erni, Peter Loretan, and Oscar Nierstrasz. Software cartography: thematic software visualization with consistent layout. *Journal of Software Maintenance and Evolution*, 22(3):191–210, April 2010.
- [13] Raula Gaikovina Kula, Daniel M. German, Ali Ouni, Takashi Ishio, and Katsuro Inoue. Do developers update their library dependencies? *Empirical Softw. Engg.*, 23(1):384–417, feb 2018. doi:10.1007/s10664-017-9521-5.
- [14] Patrick Lam, Jens Dietrich, and David J. Pearce. Putting the semantics into semantic versioning. In *Onward! Essays*, 2020.
- [15] Benjamin Livshits, Manu Sridharan, Yannis Smaragdakis, Ondřej Lhoták, J. Nelson Amaral, Bor-Yuh Evan Chang, Samuel Z. Guyer, Uday P. Khedker, Anders Møller, and Dimitrios Vardoulakis. In defense of soundness: A manifesto. *Commun. ACM*, 58(2):44–46, jan 2015. doi:10.1145/2644805.
- [16] Leonel Merino, Mohammad Ghafari, Craig Anslow, and Oscar Nierstrasz. A systematic literature review of software visualization evaluation. *Journal of Systems and Software*, 2018.
- [17] Alex Mullans. Keep all your packages up to date with dependabot. GitHub blog: <https://github.blog/2020-06-01-keep-all-your-packages-up-to-date-with-dependabot/>, June 2020.
- [18] Brad A. Myers and Jeffrey Stylos. Improving API usability. *Commun. ACM*, 59(6):62–69, May 2016. doi:10.1145/2896587.
- [19] Steven Raemaekers, Arie Van Deursen, and Joost Visser. Semantic versioning versus breaking changes: A study of the maven repository. In *SCAM*, pages 215–224. IEEE, 2014.
- [20] César Soto-Valero, Nicolas Harrand, Martin Monperrus, and Benoit Baudry. A comprehensive study of bloated dependencies in the Maven ecosystem. *Empirical Software Engineering*, 26(3):1–44, 2021.
- [21] Suresh Thummalapenta and Tao Xie. SpotWeb: Detecting framework hotspots and coldspots via mining open source code on the web. In *ASE*, pages 327–336. IEEE Computer Society, 2008.
- [22] Jukka Viljamaa. Reverse engineering framework reuse interfaces. In *FSE*, page 217–226, New York, NY, USA, 2003. Association for Computing Machinery. doi:10.1145/940071.940101.
- [23] Hao Zhong and Hong Mei. An empirical study on API usages. *IEEE Transactions on Software Engineering*, 45(4):319–334, December 2019.